

Restricted Boltzmann Machines

Project 3 - Math-6373 - Prof. Azencott

Simon Stolarczyk - PSID: 1201831

Adrian Radillo - PSID: 1328335

April 3, 2017

GitHub: The Python code used for this project is available at the public repository, <https://github.com/Spstolar/BMachine.git>

1 Generic Boltzmann Machines

1.1 Gibbs sampler for a Boltzmann machine

For a given configuration \mathbf{X} of the machine, let $X(i), X(j) \in \{-1, 1\}$ denote the states of nodes i and j respectively, and $W_{ij} = W_{ji}$ the weight between them ($i, j \in \{1, \dots, 30\}$).

Let σ be the sigmoid function and define

$$v_i = \sum_{j \neq i} W_{ij} X(j)$$

The Boltzmann machine stochastic dynamics by a Gibbs sampler corresponds to algorithm 1 below.

Algorithm 1 Gibbs sampler

```
1: totSweeps  $\leftarrow n$  # Set the total number of sweeps
2: for sweep in  $\{1, \dots, \text{totSweeps}\}$  do
3:   Select  $\alpha$  to be a random permutation of the nodes indices  $\{1, 2, \dots, 30\}$ 
4:   for  $i$  in  $\alpha$  do # loop over nodes
5:      $X(i) \leftarrow 1$  with probability  $\sigma(v_i)$  and  $-1$  otherwise # refresh state
6:   end for
7: end for
```

1.2 Results of generic BM

In appendix A.3, we include the code written for creating and running a BM according to the Gibbs sampler. We run the machine twice, once with random initial weights, and a second time with these same weights scaled. These

two Boltzmann machines were simulated until stabilization, after which we further simulated them for an extra 200 sweeps in order to compute their energy histograms.

1.2.1 Stabilization times

The stabilization times for both sets of initial weights were very similar: $S = 112$ sweeps for the original experiment and $S = 111$ sweeps when the initial weights were rescaled by $\frac{1}{10}$.

In order to visualize the evolution of the empirical means of each node in the machines, for each node j , and for each time step $t < S$ preceding stabilization, we color-plotted in figure 1, the boolean condition:

$$|M_t(j) - M_{t-1}(j)| < \text{THR.}$$

In figure 1, the purple color signifies that the condition is **False** and yellow that it is **True**. We observe very comparable stabilization times for both machines, but a great discrepancy in the time course of the stabilization state of the nodes. The machine with larger weights displays a greater proportions of stabilized nodes from the first sweep. On the other hand, the machine with smaller weights spends most of the first 100 sweeps with all its nodes ‘un-stabilized’. We are not sure of how to explain this phenomenon.

[\[A: Anything clever to add?\]](#)

1.2.2 Energy of stabilized BM

The Boltzmann energy of a configuration $\mathbf{X} = (X(1), \dots, X(m))$ is defined by:

$$E(\mathbf{X}) := - \sum_{i < j} W_{ij} X(i) X(j) - \sum_{i=1}^m b_i X(i), \quad (1)$$

where the first sum is over all pairs of distinct nodes and the $b_i \in \mathbb{R}$ are the thresholds associated to each node.

In figure 2, we present the histograms of the 200 energies, computed over the 200 sweeps following the stabilization time, S , of each machine. We observe that the distribution of energies is narrower and with mean closer to zero, when the initial weights are rescaled by $\frac{1}{10}$. [\[A: Explanation?\]](#)

2 Database

We chose the automatic hand-written digits classification task with the MNIST database. Each image is composed of $28 \times 28 = 784$ pixels. In the original dataset, each pixel is encoded as an 8-bit binary word, which codes for an integer value between 0 and 255, representing the gray scale intensity.

Given an image, each one of these 784 intensity values constitutes a feature, and each image belongs to one of the ten classes $\{0, 1, \dots, 9\}$.

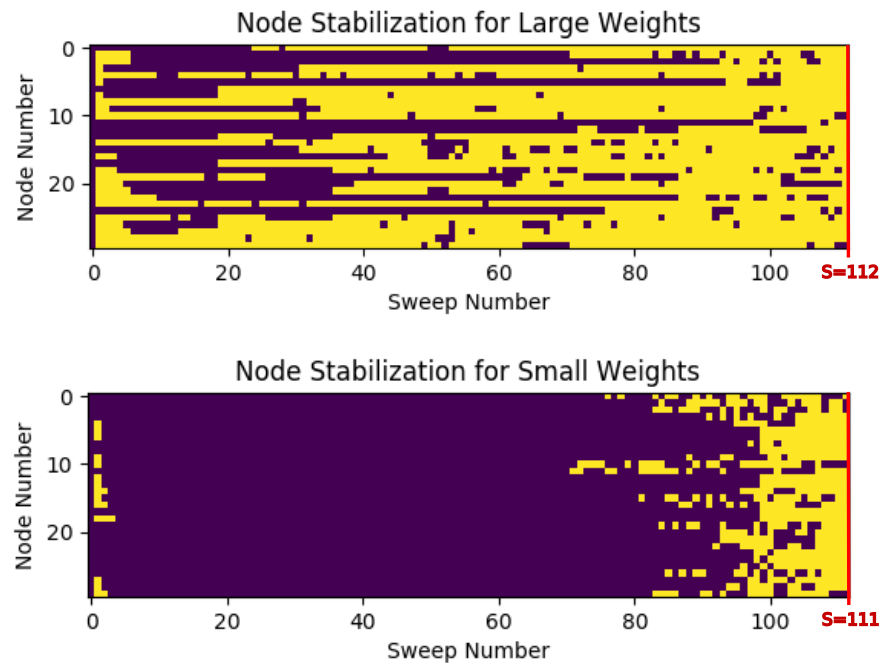


Figure 1: Time course of the stabilization status of each node during training, for the large- (top plot) and small-weight (bottom plot) cases. Purple color (not stabilized) means that the increment in the empirical mean is above the $\text{THR}=1\%$ threshold, whereas yellow color means that it is below.

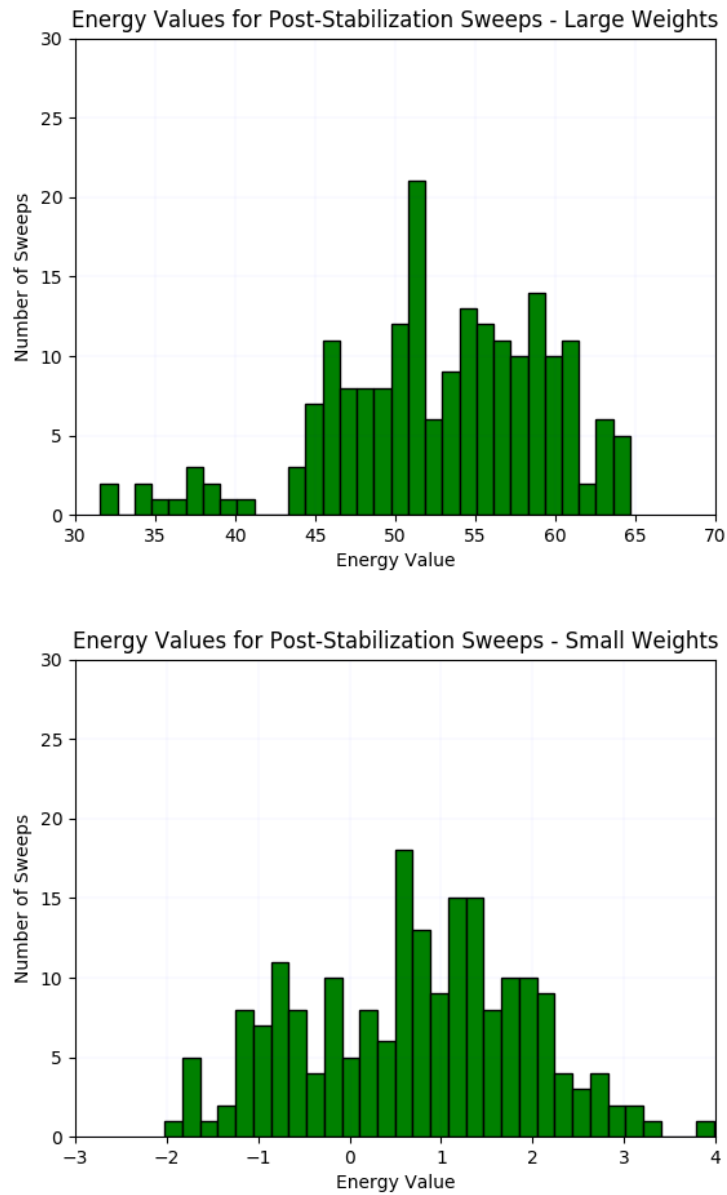


Figure 2: Histograms of energies computed over the 200 sweeps after stabilization times for large (left) and small (right) weights.

To both reduce the input size of our network and render the features interpretable by a Boltzmann machine (which traditionally is made of binary units), we followed the strategy of [1] which consists in setting all the pixels with intensity value smaller than 128 to 0 and all the remaining ones to 1. The code we wrote for this extraction and pre-processing of the data is included in appendix A.1. An example image after this ‘binarization’ was applied, is presented in figure 3. The code used to produce the image is included in appendix A.2.

Our training and test sets contain 60,000 and 10,000 cases respectively.




Figure 3: Example from the MNIST database, after undergoing our ‘binarization’ procedure.

3 Stochastic Auto-encoders based on the restricted Boltzmann machine (RBM)

After our pre-processing stage, each case from our dataset contains 1 bit of information per feature, and 784 features. It is therefore straight-forward to encode such data into the input layer of our RBM. We choose $n_3 = n_1 = 784$ and each bit from the image is mapped onto its RBM unit counterpart according to the rule: $0 \mapsto -1$ & $1 \mapsto 1$.

Our two tentative values of $h < n_1$ are 100 and 500.

3.1 Implement fast RBM learning algorithm to train the 2 auto-encoders

Our fast RBM training algorithm is summarized in Algorithm 2.

3.1.1 Batch construction

We train our RBM in the following way. First, we fix a batch size (500), a batch overlap size (100) and a total number of iterations (20) for the training. We call an iteration a presentation of the whole training set (including some redundancy¹ due to overlap between batches). Our `training()` method produces the batches and is exposed in appendix A.3.

We update the weights between each batch presentation, after having run the machine in the clamped and un-clamped mode for each example in the batch. The weight increment is based on the difference between the coactivities in both modes.

Algorithm 2 Fast RBM training

```

1: initialize RBM
2:  $batchSize \leftarrow m$                                 # batch size
3:  $batchOverlap \leftarrow \nu$                             # batch overlap size
4:  $iterations \leftarrow K$                                 # number of times training set is presented
5: for  $iter$  in  $\{1, \dots, iterations\}$  do
6:   produce batches
7:   initialize  $coactClamped$  &  $coactUnclamped$           # coactivity matrices
8:   for  $batch$  in  $batches$  do
9:      $timeStep \leftarrow$  number of batches seen since time 0
10:    for  $example$  in  $batch$  do
11:       $clamped\_run(example)$                             #1 sweep
12:      update  $coactClamped$                                 # running average
13:       $unclamped\_run(example)$                           #1 sweep
14:      update  $coactUnclamped$                             # running average
15:    end for
16:  end for
17:   $dW_{ij} \leftarrow \eta \cdot (coactClamped_{ij} - coactUnclamped_{ij})$ 
18:   $W_{ij} \leftarrow W_{ij} + dW_{ij}$ 
19: end for

```

¹Minor detail: for the last batch, if its size goes beyond the last case in the training set, we wrap the batch around the training set and append the first few examples from the training set that are needed to complete the batch.

3.2 Detailed analysis of hidden layer structure and efficiency

3.3 PCA analysis

We include the code for running principal component analysis in appendix A.4. Once we have a stable run of our machine we can easily plug in the data to get images of the hidden layer structure for the activations across inputs.

3.4 Autoencoding efficiency

Our code is quite flexible. So, once we figure out the issues with the machine we can easily tack on this classification layer and run the task.

References

- [1] Asja Fischer and Christian Igel. Training restricted Boltzmann machines: An introduction. *Pattern Recognit.*, 47(1):25–39, 2014.

Appendix A Python code

A.1 Convert MNIST binary data into numpy arrays

The following script was used to convert the image datasets into numpy arrays.

```
1 import numpy as np
2
3 # url to format of files: http://yann.lecun.com/exdb/
  mnist/
4 # definition of constants corresponding to the MNIST
  data sets
5 num_train_examples = 60000
6 num_test_examples = 10000
7 num_classes = 10
8 pixels_per_image = 784
9 bits_per_pixel = 8
10 header_bytes_images = 16
11 header_bytes_labels = 8
12
13 def bits(f):
14     bytes = (ord(b) for b in f.read()) # stores all
      bytes in the file
15     for b in bytes:
16         '''
17         The following line only returns the last bit
      of every byte
```

```
18         The yield command returns a generator as well
19         explained here:
20         http://stackoverflow.com/questions/231767/what
21         -does-the-yield-keyword-do-in-python
22         '''
23         yield (b >> 7) & 1 # I am not quite sure what
24         the & 1 is for
25
26
27 def extract_images_bin2py(num_examples, header_bytes,
28 filename, savefile, pixels_per_image=784):
29     dataset = np.zeros((num_examples, pixels_per_image
30 ))
31
32     byte_count = 1
33
34     for b in bits(open(filename, 'r')):
35
36         if byte_count <= header_bytes: # Skip header
37             bytes
38             pass
39
40         else:
41             corrected_byte_count = byte_count -
42             header_bytes - 1
43             image = corrected_byte_count /
44             pixels_per_image # Which example this
45             belongs to
46             pixel = corrected_byte_count %
47             pixels_per_image # Which pixel this
48             belongs to
49             dataset[image, pixel] = b
50
51             byte_count = byte_count + 1
52     np.save(savefile, dataset)
53
54     # arguments for training images
55     # filename='train-images-binary'
56     # savefile= 'trainSet.npy'
57
58     #arguments for test images
59     filename = 'test-images-binary'
60     savefile = 'testSet.npy'
61
62     extract_images_bin2py(num_test_examples,
63 header_bytes_images, filename, savefile)
```

The following script was used to convert the label datasets into numpy arrays.

```
1 import numpy as np
2
3
4 header_length = 64
5 num_examples = 60000
6 pixels_per_pic = 784
7 bits_per_label = 8
8 bits_per_pixel = 8
9 pic_length = pixels_per_pic * bits_per_pixel
10
11 test_label_binary = np.zeros((num_examples,
12                               bits_per_label))
12 test_label_classes = np.zeros((num_examples, 10))
13
14 def bits(f):
15     bytes = (ord(b) for b in f.read())
16     for b in bytes:
17         for i in xrange(8):
18             yield (b >> i) & 1
19
20 i = 1
21 pic_count = 0
22
23 for b in bits(open('train-labels-binary', 'r')):
24     if i <= header_length: # Preamble
25         pass
26     else:
27         example = (i - header_length - 1) /
28                   bits_per_label # Which example this
29                                   belongs to.
30         bit = ((i - header_length - 1) %
31               bits_per_label) # Which column/bit
32         test_label_binary[example, bit] = b
33
34     i += 1
35
36 for j in range(0, num_examples):
37     class_label = 0
38     for k in range(0, 8):
39         class_label += test_label_binary[j, k] * (2 **
40                                                    k)
41     test_label_classes[j, class_label] = 1 # May want
42                                             to change this to class_label.astype(int) or
```

```
        something similar.
38
39
40 np.save("train_labels.npy", test_label_classes)
```

A.2 Visualize digits from our numpy arrays

The following code was used to produce figure 3.

```
1 import matplotlib.pyplot as plt
2 import matplotlib.image as mpimg
3 import numpy as np
4
5 length = 28
6
7 dog = np.load('stabilization_small.npy')
8
9 num_plot = 0
10
11 if num_plot == 1:
12     for j in range(10):
13         img = dog[j,:].reshape((length, length)) #
14             for printing numbers
15         # img = np.random.randint(0,2,size=(length,
16             length))
17         imgplot = plt.imshow(img)
18         plt.savefig('example' + str(j) + '.png')
19 else:
20     small_plot = np.load('stabilization_small.npy')
21     img = small_plot[:112,:].T
22     imgplot = plt.imshow(img)
23     plt.title('Node Stabilization for Small Weights')
24     plt.xlabel('Sweep Number')
25     plt.ylabel('Node Number')
26     plt.savefig('stabilization_small.png')
27
28     plt.clf()
29
30     large_plot = np.load('stabilization_large.npy')
31     img = large_plot[:112, :].T
32     imgplot = plt.imshow(img)
33     plt.title('Node Stabilization for Large Weights')
34     plt.xlabel('Sweep Number')
35     plt.ylabel('Node Number')
36     plt.savefig('stabilization_large.png')
```

A.3 Create, run and train a BM

The following script defines a *BoltzmannMachine* class in Python which contains several methods and properties. The main ones are described below:

- Create an instance (`BoltzmannMachine(Nodes)`)
- Initialize weights (`create_random_weights()`)
- Run the Gibbs sampler (`run_machine()`)
- Store history of configurations (`.history`)
- Store history of energies of the system after stabilization (`.energy_history`)

The end of the script, after the class definition, performs all the computations used in section 1 and stores the energy and stabilization history in two numpy arrays.

```

1 import numpy as np
2 import time
3
4 '''
5 Parameters to modulate:
6     Learning rate = how much to change the weights by
7     for each batch
8     Batch size = how many examples to use for weight
9     change
10    Coactivity sweeps = how many sweeps to compute
11    coactivity for given weights
12    Readout sweeps = how many sweeps to run before
13    getting a readout
14    Size of training set = how many training examples
15    you'll use
16    iterations = how many times to run through the
17    training set
18 '''
19
20 def sigmoid(input_comb):
21     return 1.0 / (1 + np.exp(-input_comb))
22
23
24 def rand_bern(length):
25     # Return a random vector of -1s and 1s.
26     rand_vec = np.random.randint(0, 2, length, dtype=
27         int) # Begin with a random 0-1 draw.
28     return (rand_vec - .5) * 2 # Convert to -1, +1
29     state.
30
31

```

```
23
24 def rand_bern_with_thresh(length, fix1, fix2):
25     # Return a random vector of -1s and 1s.
26     rand_vec = np.random.randint(0, 2, length, dtype=
27         int) # Begin with a random 0-1 draw.
28     rand_vec[fix1] = 1
29     rand_vec[fix2] = 1
30     return (rand_vec - .5) * 2 # Convert to -1, +1
31     state.
32
33 def convert_binary_to_pm1(matrix):
34     """
35     Convert a 0/1 matrix to a -1/1 matrix.
36     :param matrix: A binary matrix.
37     :return: Converted matrix.
38     """
39     converted_matrix = (matrix - 0.5) * 2
40     return converted_matrix
41
42 class BoltzmannMachine(object):
43     def __init__(self, input_size, hidden_size,
44         output_size):
45         self.input_size = input_size + 1
46         self.hidden_size = hidden_size + 1
47         self.output_size = output_size
48         self.total_nodes = self.input_size + self.
49             hidden_size + self.output_size
50         self.hidden_ind = self.input_size #
51             coordinate index where the hidden layer
52             STARTS
53         self.out_ind = self.input_size + self.
54             hidden_size # coordinate index where the
55             output layer STARTS
56         self.input_thresh = self.hidden_ind - 1
57         self.hidden_thresh = self.out_ind - 1
58         self.hidden_nodes = np.arange(self.hidden_ind,
59             self.hidden_thresh)
60         self.out_nodes = np.arange(self.out_ind, self.
61             total_nodes)
62         self.clamped_visit_list = self.hidden_nodes
63         self.unclamped_visit_list = np.hstack((self.
64             hidden_nodes, self.out_nodes))
65
66         self.state = rand_bern_with_thresh(self.
```

```

        total_nodes, self.input_thresh, self.
        hidden_thresh)

58
59     self.weights = self.create_random_weights()
60     self.correct_weights()
61
62     self.in_to_hidden = np.ix_(np.arange(self.
        hidden_ind), self.hidden_nodes)
63     self.hidden_to_hidden = np.ix_(self.
        hidden_nodes, self.hidden_nodes)
64     self.hidden_to_out = np.ix_(np.arange(self.
        hidden_ind, self.out_ind), self.out_nodes)
65
66     self.batch_size = 500
67     self.inc = 400
68     self.learning_rate = .05
69     self.rate = self.learning_rate
70
71     self.history = self.state
72     self.sweeps = 100
73     self.stabilization = np.zeros((self.sweeps,
        self.total_nodes))
74     self.threshold = .01
75     self.energy_history = np.zeros(200)
76
77     def print_current_state(self):
78         print self.state
79
80     def state_energy(self):
81         """
82         Computes the current energy of the system.
83         :return: Current system energy.
84         """
85         agreement_matrix = np.outer(self.state, self.
            state) # The (i,j) entry is 1 if i,j agree
            , else -1
86         energy_contributions = agreement_matrix * self
            .weights # Element-wise product.
87         energy = 0.5 * np.sum(energy_contributions) #
            Leaving off bias.
88         return energy
89
90     def state_prob(self):
91         """
92         The (non-normalized) probability of this
            configuration. Does the whole calculation

```

```

        rather than just over some
        affected subsets.
93     :return: conditional probability of this
94     """
95
96     return np.exp(-self.state_energy())
97
98     def conditional_prob(self, node):
99         lin_sum_neighbors = np.dot(self.weights[node,
100             :], self.state)
101         return sigmoid(lin_sum_neighbors)
102
103     def update(self, node):
104         """
105         Probabilistically update a single node fixing
106         all others.
107         :param node: The number of the node to update.
108         :return: Change the state of the node
109                 according to the probabilities of the two
110                 alternatives.
111         """
112         plus_prob = self.conditional_prob(node) # P
113             ( x_j = 1 | all other node states)
114         coin_flip = np.random.binomial(1, plus_prob)
115         result = 2*(coin_flip - .5) # Convert biased
116             coin flip to -1 or 1.
117         self.state[node] = result
118
119     def simultaneous_update(self, mle=0):
120         """
121         Update ALL nodes at once. This currently will
122         update everything, including the input/
123         output nodes.
124         :param mle: Whether to pick the most likely
125                     state of each node (1), or to use the
126                     probabilistic update (0).
127         :return: Updates the state vector.
128         """
129         new_state = self.state
130         if mle == 1:
131             for node in range(0, self.total_nodes):
132                 new_state[node] = self.mle_update(node)
133             self.state = new_state
134         elif mle == 0:
135             for node in range(0, self.total_nodes):
136                 plus_prob = self.conditional_prob(node)

```

```

        ) # P( x_j = 1 | all other node
          states)
127         coin_flip = np.random.binomial(1,
          plus_prob)
128         result = 2 * (coin_flip - .5) #
          Convert biased coin flip to -1 or
          1.
129         new_state[node] = result
130         self.state = new_state
131
132     def mle_update(self, node, alter=1):
133         """
134         Update a node to it's highest probability
          state.
135         :param node: Which node to update.
136         :param alter: Actually change the node (1), or
          just return the most likely state (0).
137         :return: The most likely state for that given
          node.
138         """
139         if alter == 1:
140             if self.conditional_prob(node) > .5:
141                 self.state[node] = 1
142             else:
143                 self.state[node] = -1
144         elif alter == 0:
145             if self.conditional_prob(node) > .5:
146                 return 1
147             else:
148                 return -1
149
150     def run_machine(self, sweep_num, stabilized=0):
151         """
152         For sweep_num passes or until it stabilizes,
          update each of the nodes, except the inputs
          and thresholds.
153         :param sweep_num: A maximum number of times to
          update each node.
154         :param stabilized: Update the machine until it
          stabilizes (0), or just run the machine
          the given amount (1).
155         :return:
156         """
157         for sweep in range(sweep_num):
158             np.random.shuffle(self.
                unclamped_visit_list)

```

```
159         for node in self.unclamped_visit_list:
160             self.update(node)
161         if stabilized == 0:
162             if self.stabilization_check(sweep) ==
163                 1:
164                 break
165             # if stabilized == 1:
166             #     self.history = np.vstack((self.
167                 history, self.state))
168             #     self.energy_history[sweep] = self.
169                 state_energy()
170
171     def stabilization_check(self, sweep):
172         """
173         Check to see if the machine has stabilized by
174         this sweep.
175         :param sweep: Which sweep of updates this is
176         checking after.
177         :return: Whether or not stabilization of the
178         mean activation has occurred for 90% or
179         more of the nodes.
180         """
181         prev_mean = self.empirical_mean()
182         self.history = np.vstack((self.history, self.
183             state))
184         current_mean = self.empirical_mean()
185         difference = np.abs(current_mean - prev_mean)
186         self.stabilization[sweep, :] = np.less(
187             difference, self.threshold)
188         minimum_stabilized = np.floor(self.total_nodes
189             * .9)
190         if (np.sum(self.stabilization[sweep, :]) >
191             minimum_stabilized) & (sweep > 100):
192             # print sweep
193             # print self.stabilization[sweep, :]
194             return 1
195         else:
196             return 0
197
198     def create_random_weights(self):
199         weights = np.random.uniform(-1, 1, size=(self.
200             total_nodes, self.total_nodes)) # Random
201             weights ~ U([-1,1])
202         weights = np.triu(weights, k=1) # discard
203             lower diagonal terms (and the diagonal to
204             avoid self-connections)
```



```

190         weights = weights + weights.T # make the
           weights symmetric
191     return weights
192
193     def correct_weights(self):
194         self.weights[:self.hidden_ind, :self.
           hidden_ind] = 0 # forbids connections INP
           <-> INP
195         self.weights[-self.output_size:, -self.
           output_size:] = 0 # forbids connections
           OUT <-> OUT
196         self.weights[-self.output_size:, :self.
           input_size] = 0 # forbids connections IN
           -> OUT
197         self.weights[:self.input_size, -self.
           output_size:] = 0 # forbids connections IN
           <- OUT
198         self.weights[self.hidden_thresh, :self.out_ind
           ] = 0 # forbid hidden_thresh -> hidden &
           in
199         self.weights[:self.out_ind, self.hidden_thresh
           ] = 0 # forbid hidden -> hidden_thresh
200         np.fill_diagonal(self.weights, 0)
201         # self.weights[self.hidden_ind:, :self.
           hidden_ind] = 0
202         # self.weights[:self.hidden_ind, self.
           hidden_ind:] = 0
203
204     def check_weights(self):
205         w = 0
206         w += np.sum(self.weights[:self.hidden_ind, :
           self.hidden_ind])
207         w += np.sum(self.weights[-self.output_size:, -
           self.output_size:])
208         w += np.sum(self.weights[-self.output_size:, :
           self.input_size])
209         w += np.sum(self.weights[:self.input_size, -
           self.output_size:])
210         w += np.sum(self.weights[self.hidden_thresh,
           self.hidden_ind:self.out_ind])
211         w += np.sum(self.weights[self.hidden_ind:self.
           out_ind, self.hidden_thresh])
212         w += np.sum(self.weights.diagonal())
213         print 'Sum of the non-connecting weights: ' +
           str(w)
214

```

```
215     def empirical_mean(self, history=0):
216         """
217         Return the average state of each node for a
218             given history.
219         :param history: A record of the state of the
220             machine over a given number of sweeps,
221             defaults to the object
222             history.
223         :return: The mean activation of each node.
224         """
225         if history == 0:
226             history = self.history
227         return np.mean(history, axis=0)
228
229     def clamped_run(self, in_state, out_state,
230         sweep_num=1):
231         """
232         Runs the machine while forcing the input nodes
233             and output nodes to stay the same.
234         :param in_state: What the input node states
235             are.
236         :param out_state: What the output node states
237             are.
238         :param sweep_num: How many times to update the
239             non-clamped nodes.
240         :return: Updates the hidden states.
241         """
242         self.state[:self.input_thresh] = in_state
243         self.state[self.out_ind:] = out_state
244         for sweep in range(sweep_num):
245             np.random.shuffle(self.clamped_visit_list)
246             for node in self.clamped_visit_list:
247                 self.update(node)
248
249     def unclamped_run(self, in_state, sweep_num=1):
250         """
251         Runs the machine while forcing the input nodes
252             to stay the same.
253         :param in_state: What the input node states
254             are.
255         :param sweep_num: How many times to update the
256             non-clamped nodes.
257         :return: Updates the hidden and output states.
258         """
259         self.state[:self.input_thresh] = in_state
260         for sweep in range(sweep_num):
```

```
250         np.random.shuffle(self.  
            unclamped_visit_list)  
251         for node in self.unclamped_visit_list:  
252             self.update(node)  
253  
254     def clamped_run_mle(self, in_state, out_state):  
255         # Update the machine using the maximum  
            likelihood states.  
256         self.state[:self.input_thresh] = in_state  
257         self.state[self.out_ind:] = out_state  
258         for node in self.clamped_visit_list:  
259             self.mle_update(node)  
260  
261     def unclamped_run_mle(self, in_state):  
262         self.state[:self.input_thresh] = in_state  
263         for node in self.clamped_visit_list:  
264             self.mle_update(node)  
265  
266     def training(self, example_set, iterations):  
267         """  
268         Go through the entire example set for the  
            given number of iterations. For each  
            iteration, you divide the set  
269         into batches and then pass each batch to the  
            batch_process method to update the weights.  
            After each iteration  
270         you update the rate at which the weights are  
            changed.  
271         :param example_set: A set of unlabelled  
            examples.  
272         :param iterations: How many times to go  
            through the set for the learning.  
273         :return: The machine's weights are updated.  
274         """  
275         # Compute how to go through the batches.  
276         batch_size = self.batch_size  
277         inc = self.inc # This is to allow overlap  
            between batches, let it be about 80% of the  
            batch size.  
278         set_size = example_set.shape[0] # How many  
            examples.  
279         batches_per_iteration = int(set_size / inc) +  
            1 # How many batches will be needed.  
280         record_mse = 0  
281         ramse = np.zeros(iterations*  
            batches_per_iteration)
```

```
282         testing = 1
283
284     for it in range(iterations):
285         np.random.permutation(example_set)
286         print "Iteration: " + str(it)
287         last_batch_ind = 0 # initialize variable
288                             # for row index of last batch
289         for batch_num, b in enumerate(range(0,
290             set_size - inc, inc)):
291             batch = example_set[b:b + batch_size,
292                 :]
293             # To do learning rate decay change by
294             # batch number:
295             num_batches_seen = batch_num +
296                 batches_per_iteration * it
297             if testing == 1:
298                 print 'Working on batch ' + str(
299                     num_batches_seen)
300                 print 'Threshold units are at ' +
301                     str(self.state[self.
302                         input_thresh]) + ' and ' + str(
303                         self.state[self.hidden_thresh])
304                 print 'Current weights: ' + str(
305                     self.weights)
306                 print 'Current sum of weights: ' +
307                     str(np.sum(self.weights))
308                 self.check_weights()
309
310             self.rate = self.learning_rate / (
311                 num_batches_seen + 1)
312             self.batch_process(batch)
313             if record_mse == 1:
314                 ramse[num_batches_seen] = self.
315                     average_rmse(batch) # Compute
316                     the batch root mean square
317                     error.
318
319             last_batch_ind = b
320
321     # Manually calculate last batch. It
322     # includes some of the first and some of
323     # the last examples.
324     last_batch_ind += inc
325     wrap_around_ind = batch_size - (set_size -
326         last_batch_ind)
327     batch = np.vstack((example_set[
328         last_batch_ind:, :], example_set[
```

```

        wrap_around_ind, :]))
    self.batch_process(batch)
    if record_mse == 1:
        ramse[batches_per_iteration * (it + 1)
              -1] = self.average_rmse(batch)

    # self.rate = self.learning_rate / (1.0 +
    # it) # If we want to decrease the rate
    # more slowly.
    if record_mse == 1:
        print ramse
        np.save('root_avg_mse.npy', ramse)

def batch_process(self, batch):
    """
    Take in a batch of examples without class
    labels to be fed into autoencoder training.
    Go through each example in the batch, compute
    coactivity for clamped and unclamped runs
    adding to a running
    total that you then average to change the
    weights.
    :param batch: a matrix whose rows are examples
    .
    :return: Changes the weights of the machine.
    """
    batch_size = self.batch_size
    batch_coactivity_clamped = np.zeros((self.
        total_nodes, self.total_nodes))
    batch_coactivity_unclamped = np.zeros((self.
        total_nodes, self.total_nodes))
    for ex in range(batch_size):
        # First clamp down the input nodes and
        # output nodes and compute coactivity.
        self.state = rand_bern_with_thresh(self.
            total_nodes, self.input_thresh, self.
            hidden_thresh)
        self.clamped_run(batch[ex, :], batch[ex,
            :])
        batch_coactivity_clamped += self.
            coactivity(clamped=1, sweeps=1)
        # Next clamp down just the input nodes and
        # compute coactivity.
        self.state = rand_bern_with_thresh(self.
            total_nodes, self.input_thresh, self.
            hidden_thresh)

```

```

336         self.unclamped_run(batch[ex, :])
337         batch_coactivity_unclamped += self.
            coactivity(clamped=0, sweeps=1)
338     dw = (batch_coactivity_clamped -
            batch_coactivity_unclamped) / batch_size
339     self.weights += self.rate * dw # Not sure if
            this should be minus. TODO: find the
            correct rule.
340     self.correct_weights() # Lazy correction.
341
342     def coactivity(self, clamped=1, sweeps=1):
343         """
344         Computed the coactivity of each node pair
            averaged over a given number of sweeps.
345         :param clamped: If 0, then do unclamped
            updating. Otherwise, clamp output nodes.
346         :param sweeps: how many times to do a full
            update and compute the coactivities.
347         :return: A matrix of coactivity.
348         """
349         coactivity_matrix = np.zeros((self.total_nodes
            , self.total_nodes))
350         c_in_to_hidden = np.zeros((self.input_size,
            self.hidden_size - 1)) # minus 1 since
            threshold doesn't connect
351         c_hidden = np.zeros((self.hidden_size - 1,
            self.hidden_size - 1)) # minus 1 since
            threshold doesn't connect
352         c_hidden_to_out = np.zeros((self.hidden_size,
            self.output_size))
353
354         for s in range(0, sweeps):
355             c_in_to_hidden += np.outer(self.state[:
                self.hidden_ind], self.state[self.
                hidden_nodes])
356             c_hidden += np.outer(self.state[self.
                hidden_nodes], self.state[self.
                hidden_nodes])
357             c_hidden_to_out += np.outer(self.state[
                self.hidden_ind:self.out_ind], self.
                state[self.out_nodes])
358
359             if clamped == 0 & sweeps > 1:
360                 self.unclamped_run(self.state[:self.
                    hidden_ind])
361             elif clamped == 1 & sweeps > 1:

```

```

362         self.clamped_run(self.state[:self.
363                             hidden_ind], self.state[self.
364                             out_ind:])
365
366     coactivity_matrix[self.in_to_hidden] =
367         c_in_to_hidden
368     coactivity_matrix[self.hidden_to_hidden] =
369         c_hidden
370     coactivity_matrix[self.hidden_to_out] =
371         c_hidden_to_out
372     coactivity_matrix = (coactivity_matrix +
373                         coactivity_matrix.T) / 2.0
374     return coactivity_matrix / sweeps # TODO:
375         Check how to compute coactivity during
376         training.
377
378 def read_output(self, input_state, print_out=1):
379     # Need to fix an input and then run the
380     # machine till it has stabilized.
381     # Once stabilized, we can return both the
382     # maximizer state as well as the
383     # averages for 100 or so states.
384     sweep_num = 1000
385     self.unclamped_run(input_state, sweep_num)
386     output = np.zeros(self.output_size, dtype=
387                       float)
388     post_stab_sweeps = 100
389     for i in range(post_stab_sweeps):
390         self.unclamped_run(input_state)
391         output += self.state[self.out_ind:]
392     average_output = output / float(
393         post_stab_sweeps)
394     output_state = np.sign(average_output)
395     if print_out == 1:
396         # print (output_state == input_state)
397         print average_output
398         print np.sum(np.equal(output_state,
399                               input_state))
400     return output_state # TODO: Decide on the
401         exact rule for reading off the state.
402
403 def average_rmse(self, example_set):
404     """
405     Computes the difference between computed
406     output and input averaged over the examples
407     .

```

```
392         :param example_set:
393         :return: The square root of the average mean
394                square error.
395         """
396         num_ex = example_set.shape[0]
397         error = np.zeros(example_set.shape[1])
398         for i in range(num_ex):
399             input_state = example_set[i,:]
400             error += np.abs(input_state - self.
401                             read_output(input_state, 0))
402         total_error = np.sum(error)
403         return np.sqrt(total_error / float(num_ex))
404
405     # def modulate_params(self):
406
407 def main():
408     start_time = time.time()
409
410     # examples = np.load('toy_example_set.npy')
411     # np.random.permutation(examples)
412
413     examples = np.load('testSetSimple.npy')
414     examples = convert_binary_to_pm1(examples)
415
416     input_size = examples.shape[1]
417
418     # input_size = 3
419
420     BM = BoltzmannMachine(input_size, 30, input_size)
421
422     # BM.weights = np.ones((11,11))
423     # print BM.weights
424     # BM.correct_weights()
425     # print BM.weights
426
427     BM.run_machine(BM.sweeps)
428     BM.training(examples, 3)
429
430     ones_vec = np.ones(5)
431     neg_ones_vec = -np.ones(5)
432
433     vec_1 = np.hstack((ones_vec, neg_ones_vec))
434     vec_2 = np.hstack((ones_vec, ones_vec))
435     vec_3 = np.hstack((neg_ones_vec, ones_vec))
436     vec_4 = np.hstack((neg_ones_vec, neg_ones_vec))
```



```

436
437     print BM.read_output(vec_1)
438     print BM.read_output(vec_2)
439     print BM.read_output(vec_3)
440     print BM.read_output(vec_4)
441
442     print 'Random vectors: '
443     score = 0
444     for r in range(10):
445         rand = rand_bern(10)
446         output_state = BM.read_output(rand)
447         score += np.sum(np.equal(output_state, rand))
448         print 'In: ' + str(rand) + 'Out: ' + str(
449             output_state)
450     print str(score) + 'out of 100'
451
452     np.save('trained_weights.npy', BM.weights)
453
454     end_time = time.time()
455
456     print end_time - start_time
457
458 if __name__ == "__main__":
459     main()

```

The following code was used to make the histograms in figure 2. It opens an energy .npy file produced by the previous script and plots the result.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 x = np.load('energy_small_p1.npy')
5
6 n, bins, patches = plt.hist(x, 31, normed=0, histtype=
7     'bar', facecolor='green', alpha=1, edgecolor='black'
8     )
9
10 plt.xlabel('Energy Value')
11 plt.ylabel('Number of Sweeps')
12 plt.title('Energy Values for Post-Stabilization Sweeps
13     - Small Weights')
14
15 # plt.axis([30,70,0,30]) # This is for the large
16     weights.
17
18 plt.axis([-3,4,0,30]) # This is for the small weights
19
20

```

```
13 plt.grid(True)
14 plt.grid(color='b', linestyle='--', linewidth=.1,alpha
    =.3)
15
16 # plt.show() # If you don't want to save it, but just
    view it, use this line.
17 plt.savefig('energy_hist_small.png')
18
19 plt.clf()
20
21 x = np.load('energy_large_pl.npy')
22
23 n, bins, patches = plt.hist(x, 31, normed=0, histtype=
    'bar',facecolor='green', alpha=1, edgecolor='black'
    )
24
25 plt.xlabel('Energy Value')
26 plt.ylabel('Number of Sweeps')
27 plt.title('Energy Values for Post-Stabilization Sweeps
    - Large Weights')
28 plt.axis([30,70,0,30]) # This is for the large
    weights.
29 # plt.axis([-3,4,0,30]) # This is for the small
    weights.
30 plt.grid(True)
31 plt.grid(color='b', linestyle='--', linewidth=.1,alpha
    =.3)
32
33 # plt.show() # If you don't want to save it, but just
    view it, use this line.
34 plt.savefig('energy_hist_large.png')
```

A.4 PCA analysis

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.decomposition import PCA
4 from mpl_toolkits.mplot3d import Axes3D
5 import matplotlib.colors as colors
6
7 length_of_vecs = 100
8 my_data = np.zeros((10, length_of_vecs)) # Replace
    with data to be used.
9
```

```
10 pca = PCA(n_components=length_of_vecs) # This creates
    a PCA-doing object.
11 pca.fit(my_data)
12
13 myPCAEigs = pca.explained_variance_ # Creates the
    vector of eigenvalues for the cov matrix.
14 plt.plot(myPCAEigs,'ro') # Create a simple plot of
    the eigenvalues.
15 plt.title('Eigenvalues for Data')
16 plt.xlabel('Eigenvalue number')
17 plt.ylabel('Eigenvalue')
18 plt.savefig('eigPlot.png') # Save the plot.
19
20 plt.clf() # Clear this figure object for other use.
21
22 classes = np.identity(10) # Replace with class labels
    matrix.
23
24 class_labels = np.zeros(classes.shape[0]) # Empty
    vector to condense the class label matrix.
25 for i in range(classes.shape[0]):
26     class_labels[i] = np.argmax(classes[i,:])
27
28 num_classes = classes.shape[1]
29
30 # Generate a gradient of colors in hsv format.
31 hsv_colors = [(x*1.0/num_classes, x*0.5/num_classes, (
    num_classes-x)*0.5/num_classes) for x in range(
    num_classes)]
32 color_list = map(lambda x: colors.rgb2hex(x),
    hsv_colors) # Convert the colors to hex format.
33
34 # This makes a tuple rather than a numpy array out of
    the labels:
35 label_tuple = []
36 for i in range(classes.shape[0]):
37     label_tuple.append(int(class_labels[i]))
38
39 label_colors = np.choose(label_tuple, color_list) #
    This replaces each label with the corresponding
    color.
40
41 fig = plt.figure(1, figsize=(4,3))
42 ax = fig.add_subplot(111, projection='3d')
43 pca = PCA(n_components=3) # Now we force projection
    onto the principal components.
```

```
44 pca.fit(my_data)
45 data_proj = pca.transform(my_data)
46
47 num_pts = 10 # Let this be 100 or 1000. Too many
               points makes things too messy.
48
49 ax.scatter(data_proj[:num_pts,0],data_proj[:num_pts
        ,1], data_proj[:num_pts,2], c=label_colors[:num_pts
        ], cmap=plt.cm.spectral)
50
51 plt.axis('on')
52
53 # Remove the labels and their silly, imaginary units.
54 ax.set_xticklabels([])
55 ax.set_yticklabels([])
56 ax.set_zticklabels([])
57
58 plt.title('PCA Plot for Data')
59
60 plt.savefig('pca_example.png')
```